

# class08 Mini Project

Leah Johnson PID: A17394690

## Table of contents

Background . . . . .	1
Data import . . . . .	1
Principal Component Analysis (PCA) . . . . .	4
<b>Create a ‘biplot()’ of wisc.pr</b>	<b>5</b>
<b>Communicating PCA results</b>	<b>7</b>
Hierarchical Clustering . . . . .	8
Clustering PCA results . . . . .	10
Prediction . . . . .	13

## Background

In today's class we will apply the methods and techniques clustering and PCA to help make sense of a real world breast cancer fine needle aspiration (FNA) biopsy

## Data import

We start by importing the data. It is a csv file so we will use the ‘read.csv()’ function.

```
fna.data <- "WisconsinCancer.csv"  
  
wisc.df <- read.csv(fna.data, row.names=1)  
  
wisc.data <- read.csv(fna.data, row.names=1)
```

```
head(wisc.data)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.80	1001.0
842517	M	20.57	17.77	132.90	1326.0
84300903	M	19.69	21.25	130.00	1203.0
84348301	M	11.42	20.38	77.58	386.1
84358402	M	20.29	14.34	135.10	1297.0
843786	M	12.45	15.70	82.57	477.1
	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean	
842302	0.11840	0.27760	0.3001	0.14710	
842517	0.08474	0.07864	0.0869	0.07017	
84300903	0.10960	0.15990	0.1974	0.12790	
84348301	0.14250	0.28390	0.2414	0.10520	
84358402	0.10030	0.13280	0.1980	0.10430	
843786	0.12780	0.17000	0.1578	0.08089	
	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
842302	0.2419		0.07871	1.0950	0.9053
842517	0.1812		0.05667	0.5435	0.7339
84300903	0.2069		0.05999	0.7456	0.7869
84348301	0.2597		0.09744	0.4956	1.1560
84358402	0.1809		0.05883	0.7572	0.7813
843786	0.2087		0.07613	0.3345	0.8902
	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se
842302	153.40	0.006399	0.04904	0.05373	0.01587
842517	74.08	0.005225	0.01308	0.01860	0.01340
84300903	94.03	0.006150	0.04006	0.03832	0.02058
84348301	27.23	0.009110	0.07458	0.05661	0.01867
84358402	94.44	0.011490	0.02461	0.05688	0.01885
843786	27.19	0.007510	0.03345	0.03672	0.01137
	symmetry_se	fractal_dimension_se	radius_worst	texture_worst	
842302	0.03003		0.006193	25.38	17.33
842517	0.01389		0.003532	24.99	23.41
84300903	0.02250		0.004571	23.57	25.53
84348301	0.05963		0.009208	14.91	26.50
84358402	0.01756		0.005115	22.54	16.67
843786	0.02165		0.005082	15.47	23.75
	perimeter_worst	area_worst	smoothness_worst	compactness_worst	
842302	184.60	2019.0	0.1622	0.6656	
842517	158.80	1956.0	0.1238	0.1866	
84300903	152.50	1709.0	0.1444	0.4245	
84348301	98.87	567.7	0.2098	0.8663	

84358402	152.20	1575.0	0.1374	0.2050
843786	103.40	741.6	0.1791	0.5249
	concavity_worst	concave.points_worst	symmetry_worst	
842302	0.7119	0.2654	0.4601	
842517	0.2416	0.1860	0.2750	
84300903	0.4504	0.2430	0.3613	
84348301	0.6869	0.2575	0.6638	
84358402	0.4000	0.1625	0.2364	
843786	0.5355	0.1741	0.3985	
	fractal_dimension_worst			
842302		0.11890		
842517		0.08902		
84300903		0.08758		
84348301		0.17300		
84358402		0.07678		
843786		0.12440		

We make sure to remove the first ‘diagnosis’ column - I don’t want to use this for my machine learning models. We will use it later to compare our results to the expert diagnosis.

```
wisc.data <- wisc.df[,-1]
diagnosis <- wisc.df$diagnosis
```

Q1. How many observations are in the dataset?

```
nrow(wisc.data)
```

[1] 569

569 observations

Q2. How many observations have a malignant diagnosis?

```
table(diagnosis)
```

```
diagnosis
  B    M
357 212
```

```
sum(wisc.df$diagnosis == "M")
```

```
[1] 212
```

212 observations have a malignant diagnosis.

Q3. How many variables/features in the data are suffixed with \_mean?

```
length( grep("_mean", colnames(wisc.df)) )
```

```
[1] 10
```

10 variables in the data are suffixed with \_mean.

## Principal Component Analysis (PCA)

The main function here is ‘prcomp()’ and we want to make sure we set the optional argument ‘scale=TRUE’:

```
wisc.pr <- prcomp(wisc.data, scale=TRUE)
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997

	PC29	PC30
Standard deviation	0.02736	0.01153
Proportion of Variance	0.00002	0.00000
Cumulative Proportion	1.00000	1.00000

Q4. From your results, what proportion of the original variance is captured by the first principal component (PC1)?

PC1 captures 44.27% of the original variance.

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?

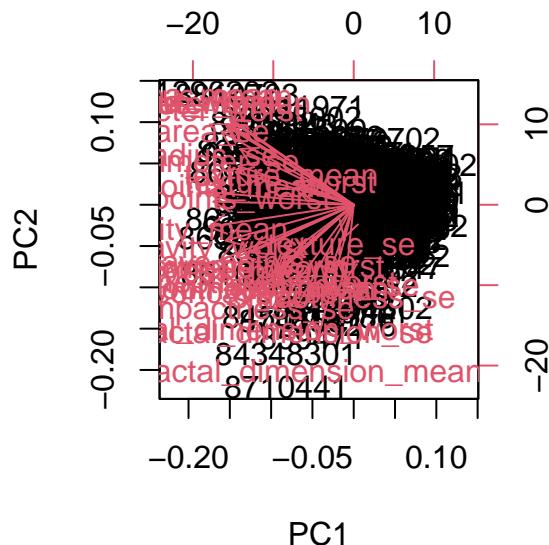
3 PCs are required to describe at least 70% of the original variance in the data.

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

7 PCs are required to describe at least 90% of the original variance in the data.

## Create a 'biplot()' of wisc.pr

```
biplot(wisc.pr)
```



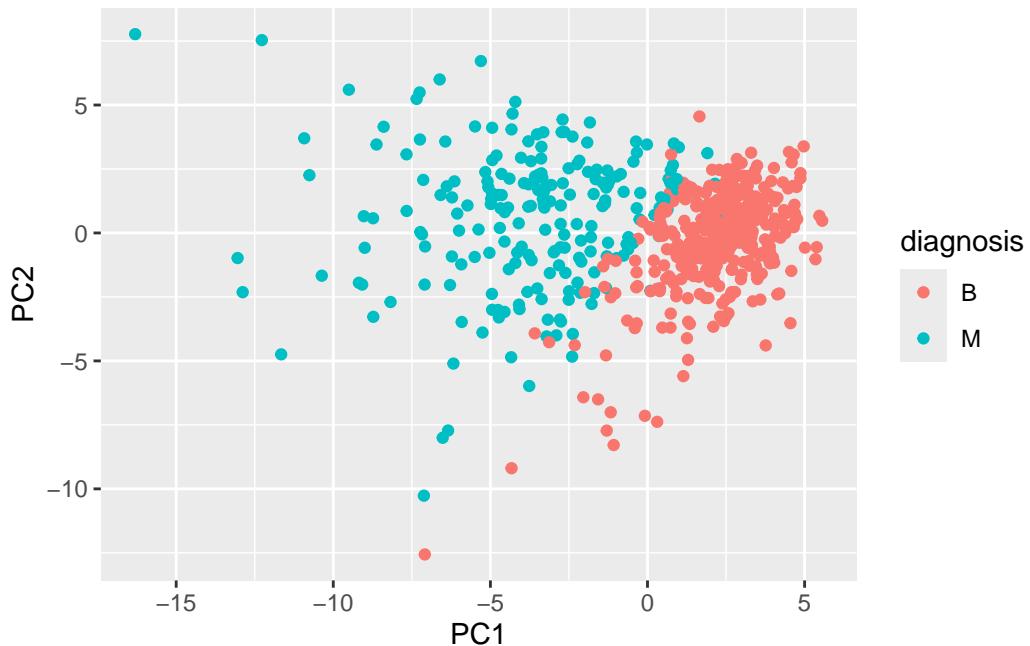
Q7. What stands out to you about this plot? Is it easy or difficult to understand?  
Why?

I notice that the row names are the labels for each of the points, therefore there is lots of overlapping text making the data difficult to visualize.

Our main PCA “score plot” or “PC plot” of results:

```
library(ggplot2)
```

```
ggplot(wisc.pr$x) +  
  aes(PC1, PC2, col=diagnosis) +  
  geom_point()
```

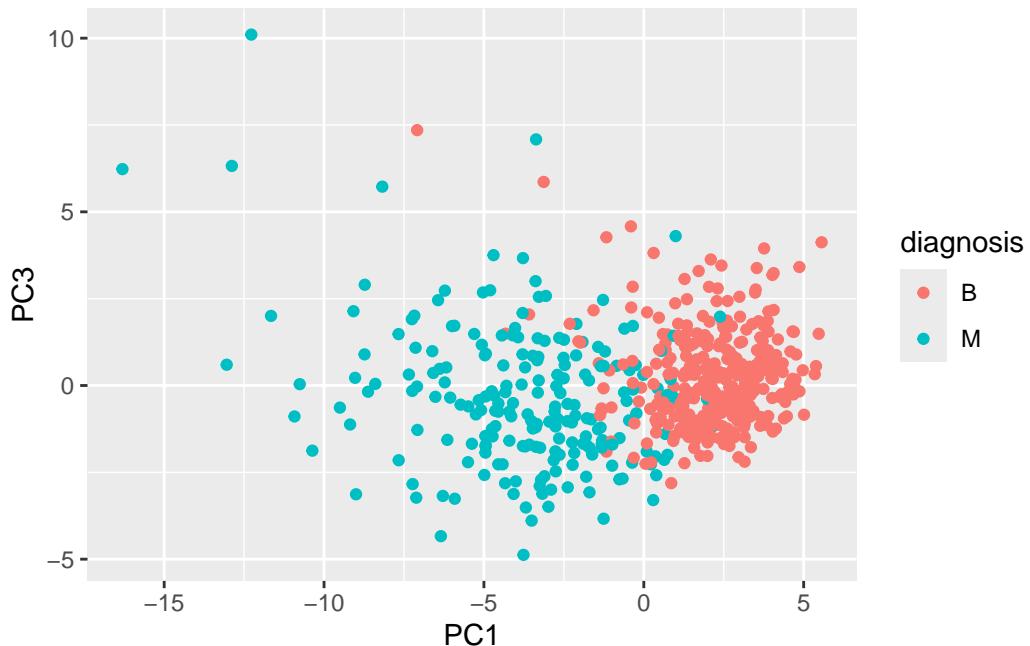


Each point represents each benign or malignant sample and its cell characteristics.

Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

I notice that majority of the points on the PC1 and PC3 plot are placed lower on the graph, the shape is very similar to the PC1 vs. PC2 graph and almost visually looks like a horizontally flipped version of the PC1 and PC2 graph.

```
ggplot(wisc.pr$x) +
  aes(PC1, PC3, col=diagnosis) +
  geom_point()
```



## Communicating PCA results

Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`? This tells us how much this original feature contributes to the first PC. Are there any features with larger contributions than this one?

`"wisc.pr$rotation[,1]"` will tell us which of the original feature measurements contribute most to PC1. There are potential features with larger contributions than this one.

```
summary(wisc.pr$rotation["concave.points_mean",1])
```

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
<code>concave.points_mean</code>	-0.2609	-0.2609	-0.2609	-0.2609	-0.2609	-0.2609

```
summary(wisc.pr$rotation["concave.points_worst",2])
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.008257	0.008257	0.008257	0.008257	0.008257	0.008257

## Hierarchical Clustering

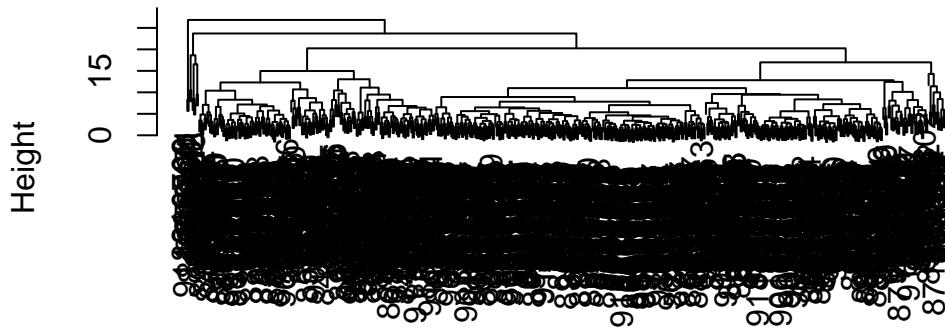
```
data.scaled <- scale(wisc.data)

data.dist <- dist(data.scaled)

wisc.hclust <- hclust(data.dist, method = "complete")

plot(wisc.hclust)
abline(wisc.data, col="red", lty=2)
```

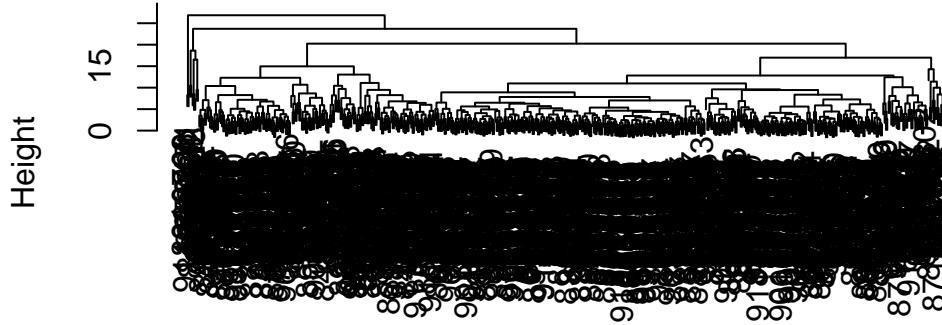
Cluster Dendrogram



```
data.dist  
hclust (*, "complete")
```

```
wisc.hclust <- hclust (dist( scale(wisc.data) ) )  
plot(wisc.hclust)
```

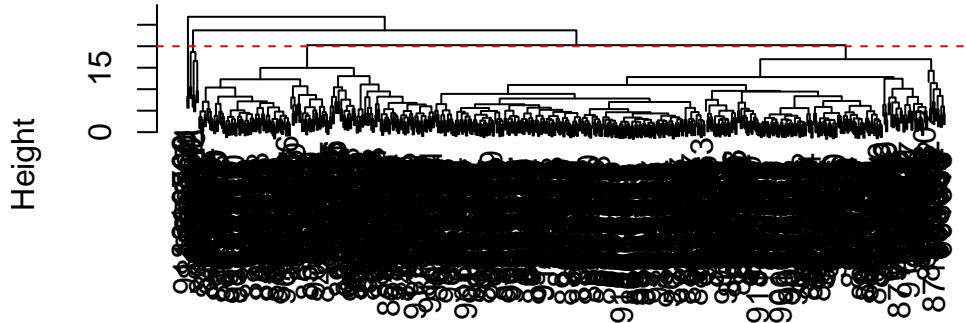
## Cluster Dendrogram



```
dist(scale(wisc.data))  
hclust (*, "complete")
```

```
plot(wisc.hclust)  
abline(h=20, col="red", lty=2)
```

## Cluster Dendrogram



```
dist(scale(wisc.data))  
hclust (*, "complete")
```

You can also use the 'cutree()' function with an argument of 'k=4' rather than 'h=height'

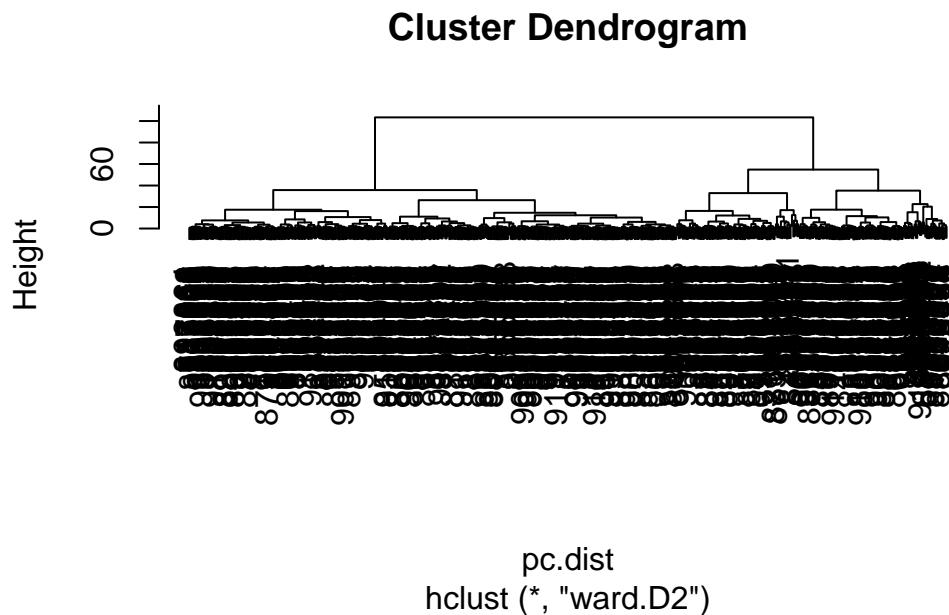
```
wisc.hclust.clusters <- cutree(wisc.hclust, k=4)
table(wisc.hclust.clusters)
```

```
wisc.hclust.clusters
 1   2   3   4
177   7 383   2
```

## Clustering PCA results

Here we will take our PCA results and use those as input for clustering. In other words our ‘wisc.pr\$x’ scores that we plotted above (the main output from PCA - how the data lie on our new principal component axis/variables) and use a subset of PCs as input for ‘hclust()’

```
pc.dist <- dist( wisc.pr$x[,1:3] )
wisc.pr.hclust <- hclust(pc.dist, method="ward.D2")
plot(wisc.pr.hclust)
```



Q12. Which method gives your favorite results for the same data.dist dataset?  
Explain your reasoning.

“ward.D2” is my favorite for obtaining the same results for the same data.dist dataset because it organizes the clusters to make the dendrogram clearer by decreasing within-group, or within-cluster, variance.

Cut the dendrogram/tree into two main groups/clusters:

```
grps <- cutree(wisc.pr.hclust, k=2)
table(grps)
```

```
grps
  1   2
203 366
```

I want to know how the clustering into ‘grps’ with values of 1 or 2 correspond the expert ‘diagnosis’

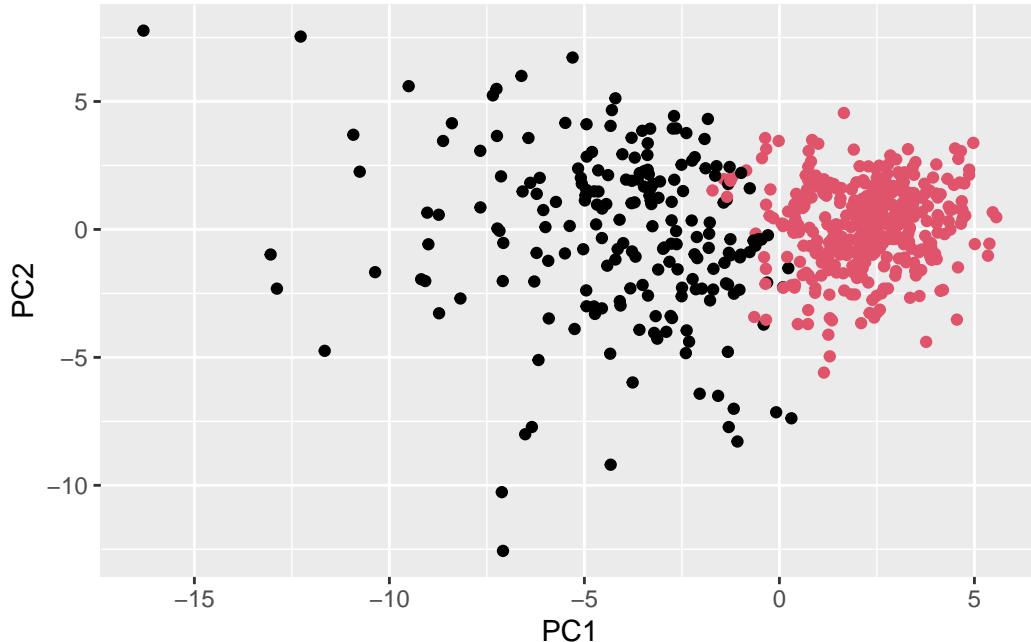
```
table(grps, diagnosis)
```

```
diagnosis
grps   B    M
  1  24 179
  2 333  33
```

My clustering **group 1** are mostly “M” diagnosis (179) and my clustering **group 2** are mostly “B” diagnosis.

24 False positives (FP) 179 True positives (TP) 333 True negatives (TN) 33 False negatives (FN)

```
ggplot(wisc.pr$x) +
  aes(PC1, PC2) +
  geom_point(col=grps)
```



Q13. How well does the newly created hclust model with two clusters separate out the two “M” and “B” diagnoses?

There is greater separation between “M” and “B” diagnoses with 2 clusters.

```
d <- dist(wisc.pr$x[, 1:7])
wisc.pr.hclust <- hclust(d, method="ward.D2")
wisc.pr.hclust.clusters <- cutree(wisc.pr.hclust, k=2)
table(wisc.pr.hclust.clusters, diagnosis)
```

	diagnosis	
wisc.pr.hclust.clusters	B	M
1	28	188
2	329	24

Q14. How well do the hierarchical clustering models you created in the previous sections (i.e. without first doing PCA) do in terms of separating the diagnoses? Again, use the table() function to compare the output of each model (wisc.hclust.clusters and wisc.pr.hclust.clusters) with the vector containing the actual diagnoses.

Prior to PCA, the separation of the diagnoses was not as concise in that there were four clusters to consider. After PCA, there were 2 clusters to display the greatest variation, or the greatest contributors to both benign and malignant tumors.

```
table(wisc.hclust.clusters, diagnosis)
```

		diagnosis	
wisc.hclust.clusters	B	M	
1	12	165	
2	2	5	
3	343	40	
4	0	2	

## Prediction

Q16. Which of these new patients should we prioritize for follow up based on your results?

Based on the new results from the new data, we should prioritize patient 2 for follow up since they fall closer to the points representing the “M” or malignancy chunk displayed in the plot.